

David Lau 253 Reuse distance report

Functionality

To test the program, I created a vector of integers: [4, 3, 4, 5, 5, 5, 6, 4]. This should return a hashmap from reuse distances to the number of times they occurred of 4 instances of infinity(None), 2 instances of 1, 1 of 3, and 1 of 2. Printing this hashmap displays the following:

```
{Some(1): 2, Some(3): 1, None: 4, Some(2): 1}
```

Showing that the program works correctly.

Complexity

The time complexity will be $O(nm)$ by the following analysis: The program will loop n times. For every loop, the program will add the current element to all the hash sets in the hashmap, which will take $O(m)$ time. So the run time is $O(n*m)$.

The space complexity will be $O(n + m^2)$. N slots will be filled by the input, and the values are stored in a hashmap of key type of the input element, into hash sets of the same element type, which has m variants, so m hash sets of size m can be stored.

Speed Tests

Top value is reuses per second, bottom is number of seconds to run total. The tests were ran with `cargo run --release`.

10^3 :

```
811366.2677719638  
0.001257873  
[dlau3@cycle1 src]$
```

10^6 :

```
701.16786572478
1426.200895315
[dlau3@cycle1 src]$
```

10^9 :

Ran for longer than an hour.

Analysis

The runtime is almost exclusively determined by the complexity. This is because essentially all the operations used are $O(1)$ with the exception of the costly `hashset.clear` function, and mostly manipulating pointers. To show this, examples of the time with a given n or m along with their resulting time, along with the same for $2n$ and $2m$, were tested and the results shown to show how they affected the time.

n, m :

```
16.7178771
```

$2n, m$:

```
34.8204955
```

$n, 2m$:

```
30.2060845
```

$2n, 2m$:

```
59.559163
```

This demonstrates the essentially proportional effect. The doubled input seems to result in a slightly larger increase than the doubled key number, which I am unaware of the cause of.